

Image classification

RAFAŁ KLINOWSKI, JAKUB CISOWSKI

BIOLOGICALLY INSPIRED ARTIFICIAL INTELLIGENCE, SEM. 6

What is image classification?

- ❖ Categorizing images
- ❖ Tries to answer the question: „What is in the picture?”
- ❖ Based on pixels or vectors within an image
- ❖ Tries to recognize patterns within an image to make predictions
- ❖ We explore „single-label classification” which only assigns one value to an image

Roses



Tulips



Dog



Cat



Real life examples of image recognition

- ❖ Medicine – eg. to help doctors detect abnormalities such as cancer
- ❖ Self-driving cars
- ❖ Face ID and other facial recognition systems
- ❖ Retail – eg. the ability to try on a piece of clothing online

Our project

Distinction between cat and dog breeds

- ❖ One step further than just saying „dog or cat”
- ❖ Input: an image
- ❖ Output: name of dog or cat breed as well as certainty (in %)

98% Beagle



84% Birman

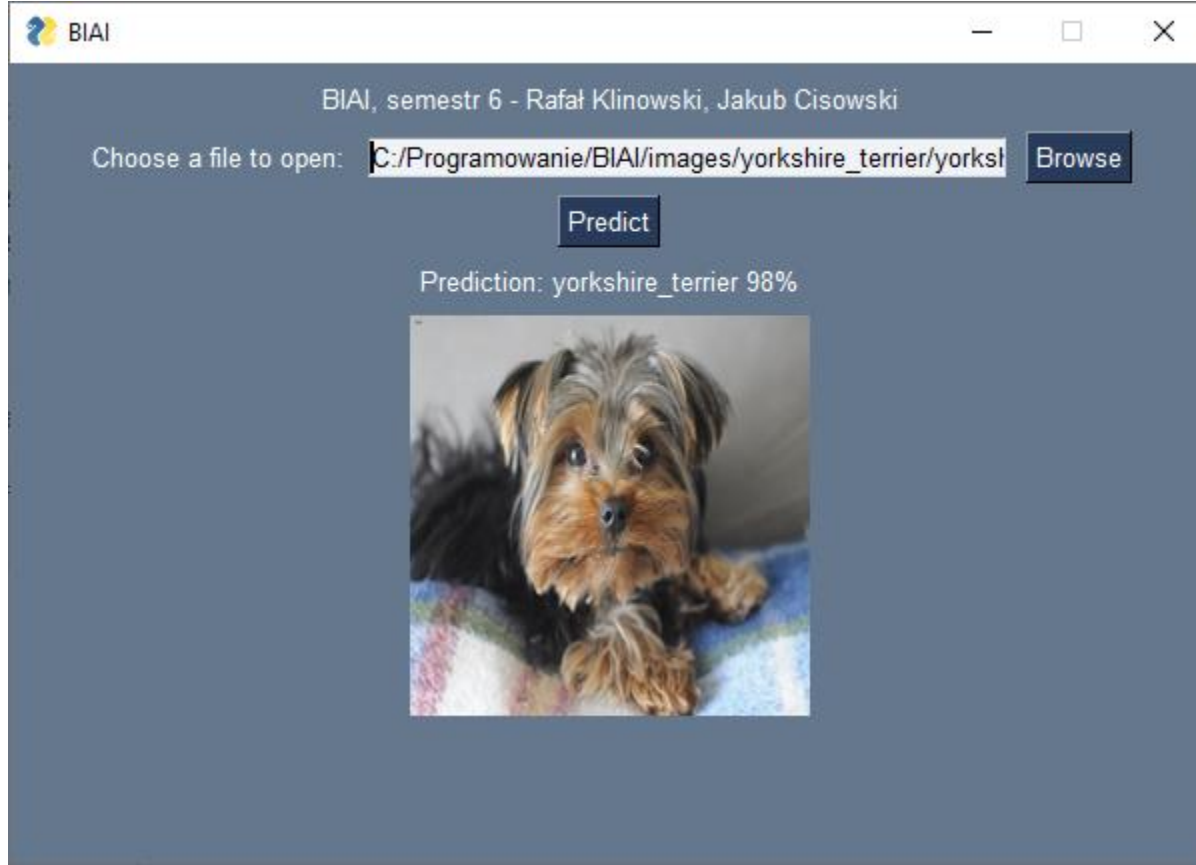


Dataset used for training

- ❖ Courtesy of The University of Oxford
- ❖ JPG image files
- ❖ 37 cat or dog races
- ❖ About 200 images for each breed
- ❖ Over 7390 images in total
- ❖ Various sizes (have to be standardized before training)

Description of technology

- ❖ Programming language: Python
- ❖ TensorFlow – used to create, train, save and load model
- ❖ Keras – used to simplify data loading and dividing into different categories
- ❖ Matplotlib – used to generate plots describing model performance, accuracy
- ❖ Algorithm:
 - Divide images into groups based on breed
 - Calculate minimum, maximum and average image size (used when creating model)
 - Load data, split into two groups: training, validation
 - Create, compile, train and evaluate model
 - Save model as a group of files



Research

We built and trained our model multiple times with different parameters to test their impact on performance

Key parts of the code

target_size, input_shape

- ❖ represents the size to which all images will be resized when first loaded
- ❖ larger values require significantly more memory and longer training time
- ❖ smaller values mean there are less trainable parameters – model will be less accurate
- ❖ common values include: 96x96, 160x160, 224x224

```
target_size = (160, 160)  
input_shape = (target_size[0], target_size[1], 3)
```


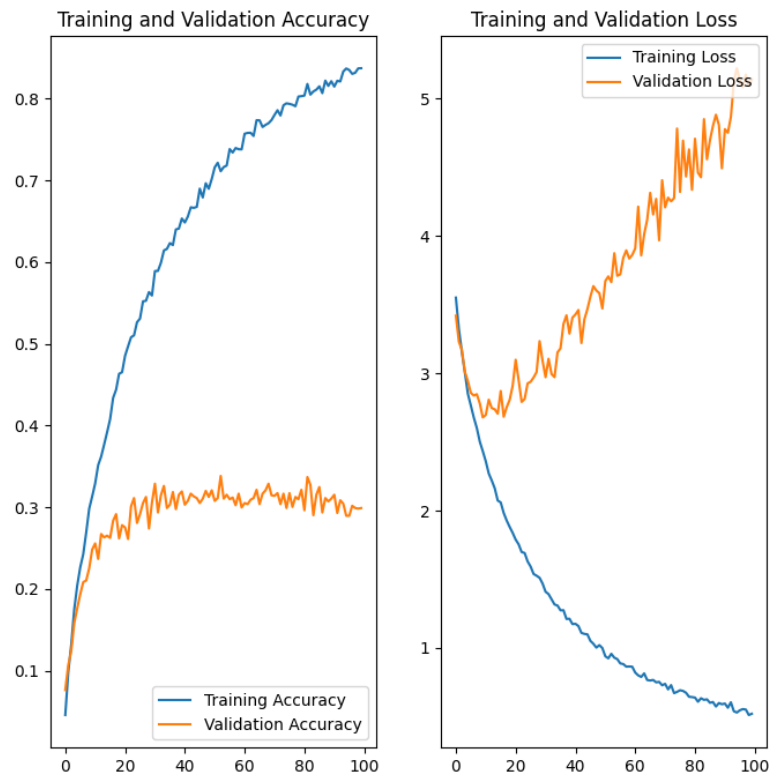
target_size=(224,224)

(note: different scale)

target_size=(96,96)

training time \approx 20 hours

training time \approx 45 minutes



batch_size

- ❖ the number of images in a batch, used to perform one step of training
- ❖ if too small, then we end up approximating the function with too few examples that might not be representative of all training data
- ❖ if too large, requires a lot of memory and can even run out of it during training which is then interrupted
- ❖ common values are 16, 32, 64

```
batch_size = 16
```

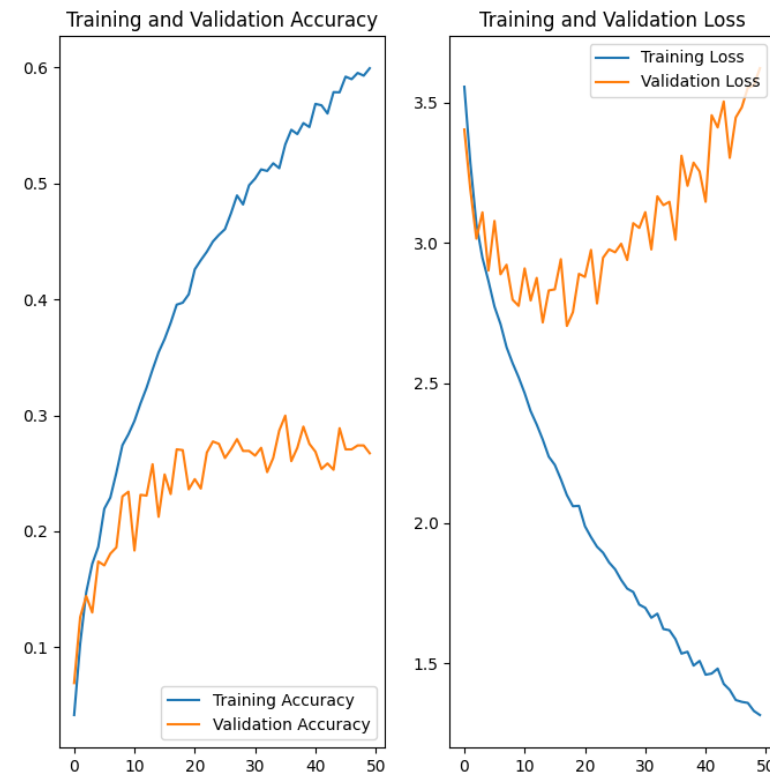
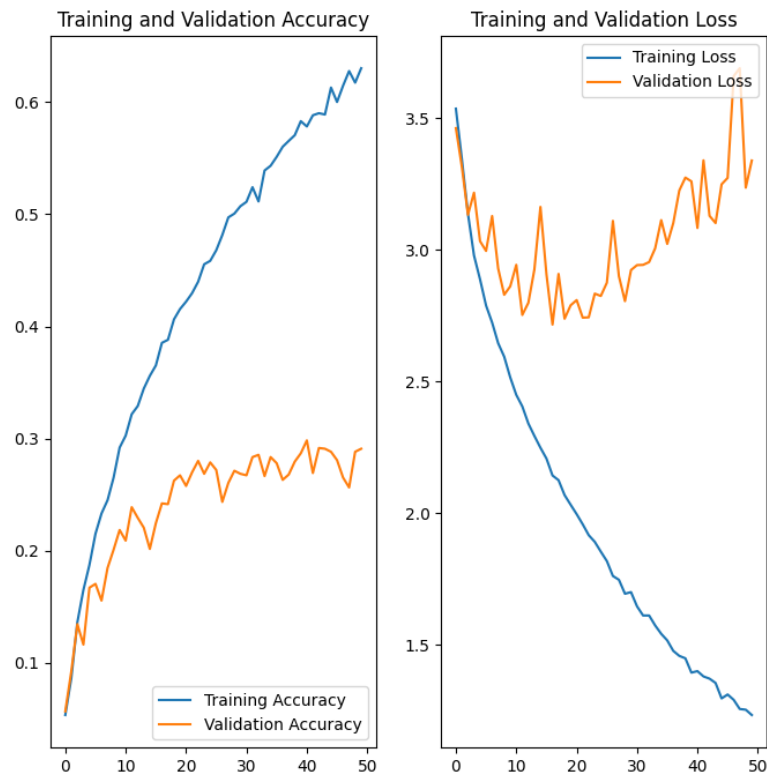
batch_size=16

(note: not enough memory to run 32)

batch_size=8

training time \approx 45 minutes

training time \approx 3 hours




validation_split

- ❖ the percentage of training images used for validation
- ❖ the rest of images is used for training only
- ❖ the dataset we used did not provide a separate set for validation – thus we had to use split
- ❖ common value is 20%

```
validation_split=0.2
```

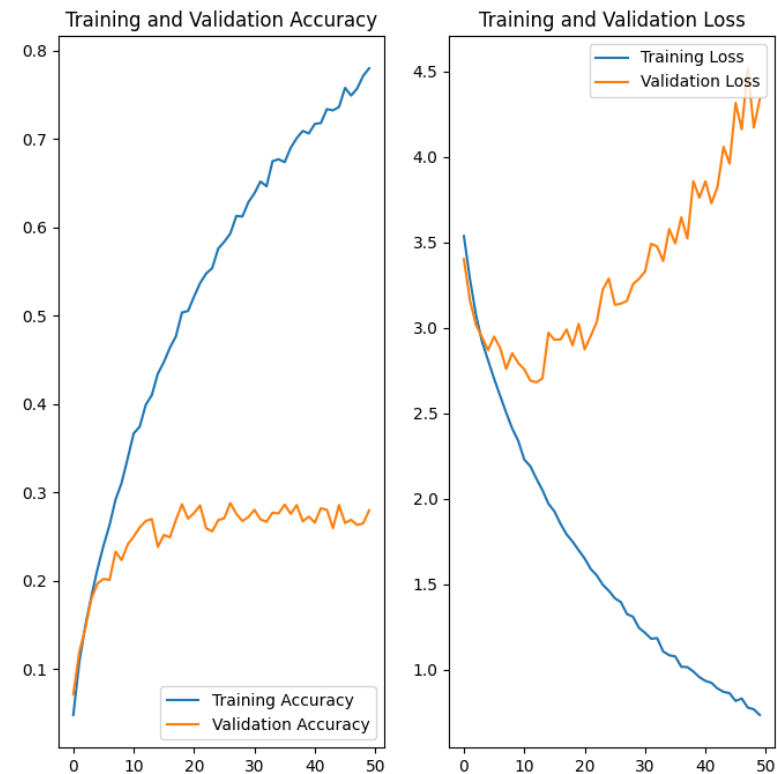
validation_split=0.2

training time \approx 45 minutes



validation_split=0.3

training time \approx 45 minutes



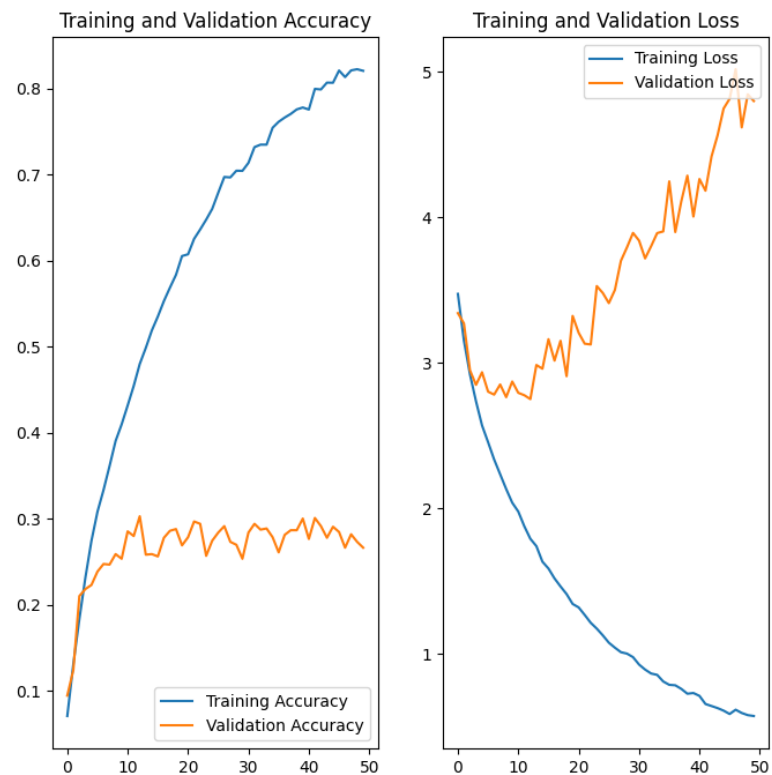
shuffle

- ❖ random ordering of images in training set
- ❖ commonly used to improve training accuracy
- ❖ prevents the model from learning the order of images
- ❖ helps with training time as well due to „prefetch” (some images are ready to be used in memory)
- ❖ shuffle parameter tells how many points to keep in memory for random drawing

```
train_ds = tf.keras.preprocessing.image_dataset_from_directory(  
    data_dir,  
    validation_split=0.2,  
    subset="training",  
    seed=6798,  
    image_size=target_size,  
    batch_size=batch_size)  
  
train_ds = train_ds.cache().shuffle(1000).prefetch(buffer_size=AUTOTUNE)
```

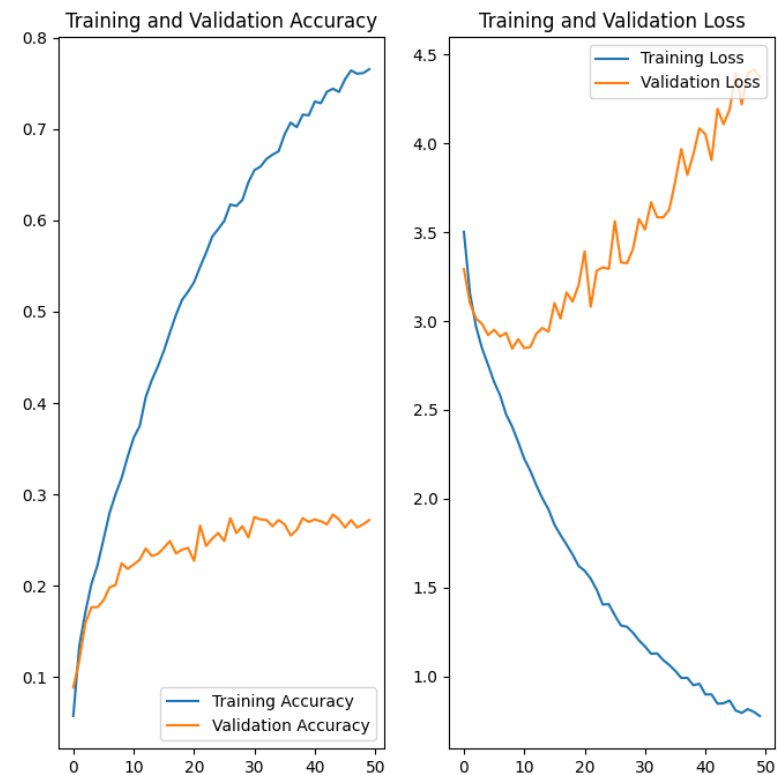
shuffle(1000)

training time \approx 45 minutes



no shuffle

training time \approx 50 minutes



data_augmentation

- ❖ helps diversify data
- ❖ improves accuracy
- ❖ we used three functions: RandomFlip (flips some of the images), RandomRotation (rotates some of the images slightly) and RandomZoom (zooms in or out on some of the images slightly)
- ❖ tested values: 0.1 and 0.2

```
data_augmentation = keras.Sequential(  
    [  
        layers.RandomFlip("horizontal",  
                           input_shape=input_shape),  
        layers.RandomRotation(0.1),  
        layers.RandomZoom(0.1),  
    ]  
)
```

Creating the model

- ❖ we could use many different layers in different orders or with different parameters
- ❖ most room for experimenting

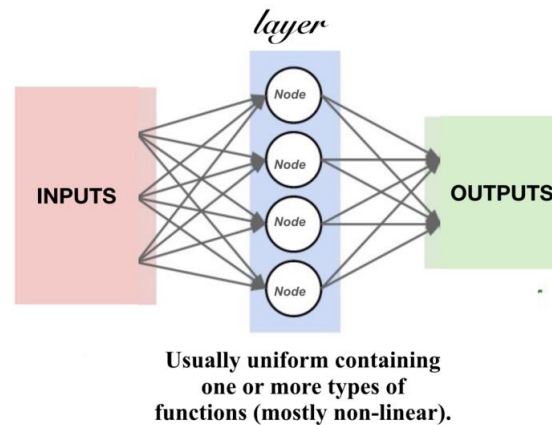
```
model = Sequential([
    data_augmentation,
    layers.Rescaling(1. / 255, input_shape=input_shape),      # Normalize colors
    layers.Conv2D(32, (3, 3), activation='relu'),
    layers.MaxPooling2D((2, 2)),
    layers.Conv2D(64, (3, 3), activation='relu'),
    layers.MaxPooling2D((2, 2)),
    layers.Conv2D(64, (3, 3), activation='relu'),
    layers.Dropout(0.2),      # Dropout 20% of the nodes to increase validation accuracy
    layers.Flatten(),
    layers.Dense(64, activation='relu'),
    layers.Dense(num_classes)
])
```

Layers used in the final model

- ❖ The model used is Sequential, which means the layers are applied in order
- ❖ data_augmentation – explained before
- ❖ Rescaling – normalizes pixel values for images (from 0 – 255 to 0.0 – 1.0)
- ❖ Conv2D – Convolution Layer, performs element-wise multiplication on some parts of the input matrix at a time, outputs a different matrix
- ❖ MaxPooling2D – downsamples the input for each channel of input by choosing max value from each feature vector patch
- ❖ Dropout – randomly drops out (sets to zero) some inputs
- ❖ Flatten – flattens the input (reduces multidimensional vectors to a single dimension array)

What is a layer?

- ❖ A layer is a structured used to pass information
- ❖ Takes one or more inputs, applies a transformation, and has one or more outputs
- ❖ Input itself is a layer too
- ❖ Transformations are mostly non-linear functions

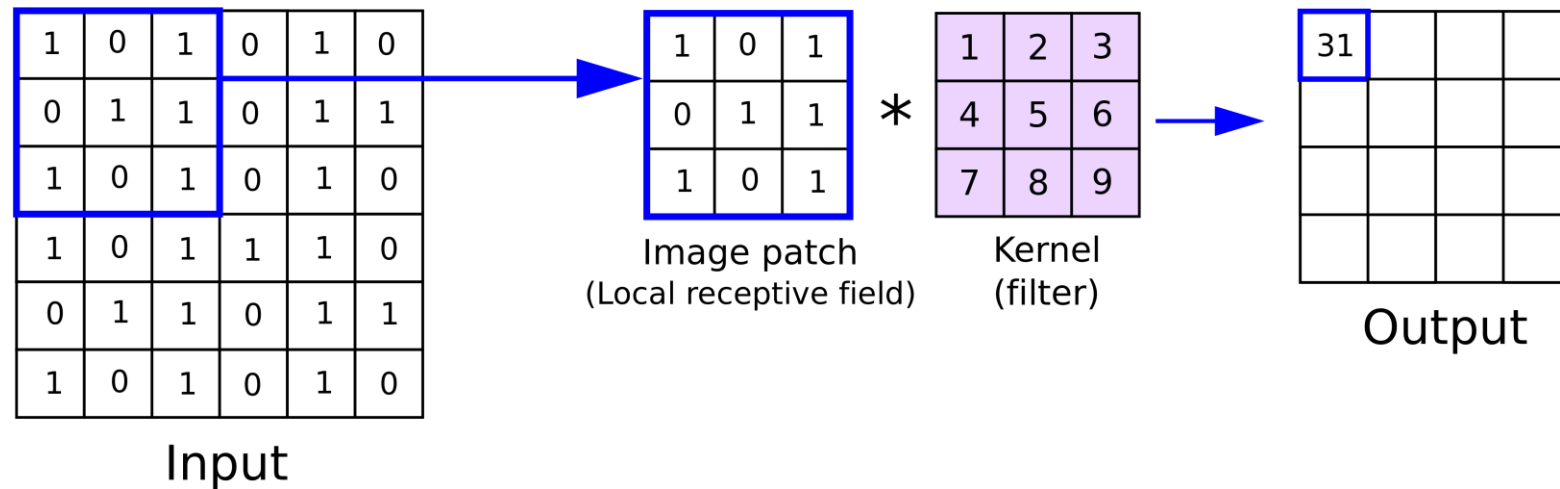


'Rescaling' layer

- ❖ It's necessary to normalize input data
- ❖ Makes all layers and features work on the same level of scale (0.0 – 1.0) instead of having different scales (0-1, 0-255, etc.)

'Conv2D' layer

- ❖ Performs element-wise multiplication on some parts of the input matrix at a time
- ❖ Outputs a different matrix

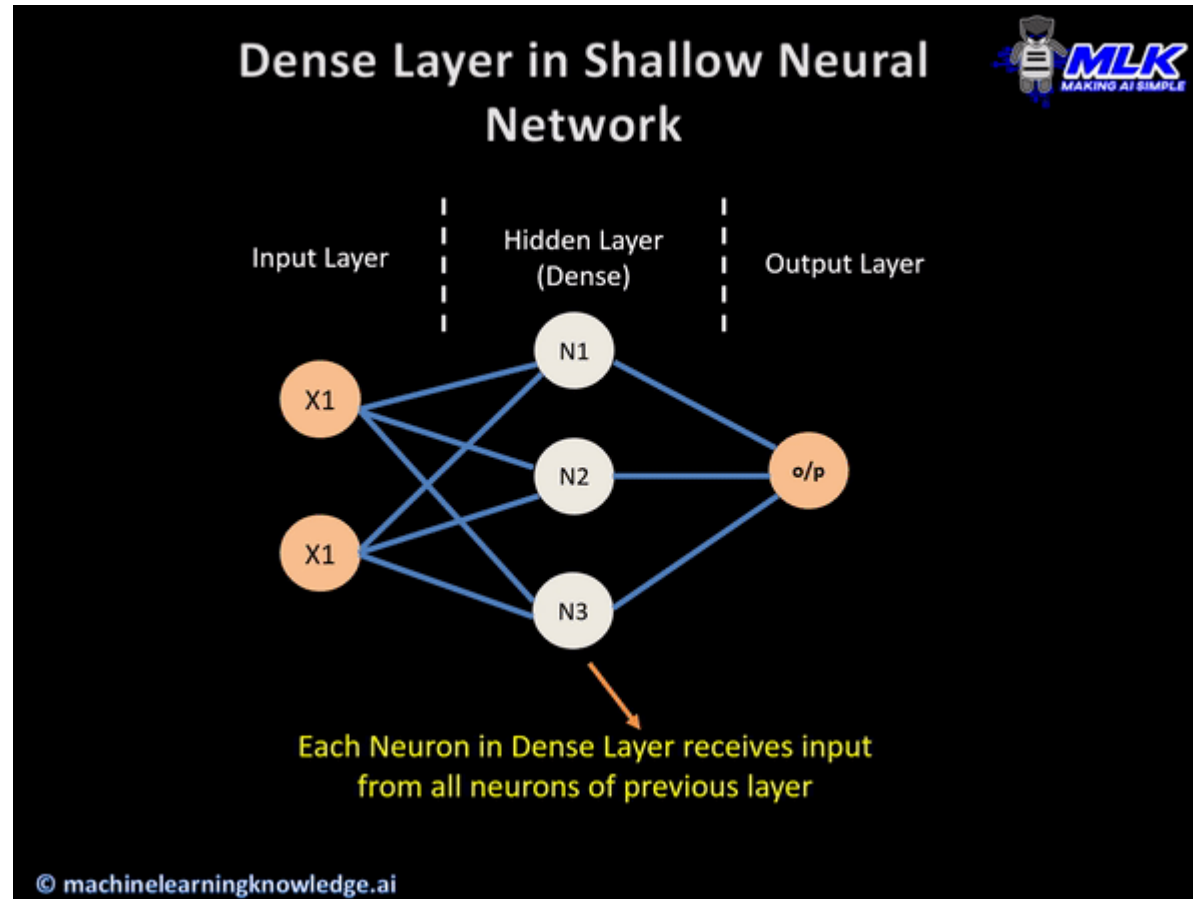


'Dense' layer

- ❖ applies a non-linear transform on data
- ❖ each neuron in Dense layer receives input from all neurons from previous layer
- ❖ activation function passed as parameter
- ❖ for activation function we used ReLU – returns input if it's positive, otherwise returns zero; it's fast, simple and universally works well
- ❖ parameter – number of outputs
- ❖ second Dense layer – output size is equal to the amount of classes we have (37)

```
layers.Dense(64, activation='relu'),  
layers.Dense(num_classes)
```

'Dense' layer – cont.




'Dropout' layer


- ❖ Used to prevent overfitting
- ❖ Overfitting – situation when model predicts „too well” on training data, and therefore is not very accurate at making predictions from outside of training set
- ❖ By randomly dropping out some of inputs we get a more diverse, less dependent set

```
layers.Dropout(0.2)
```


default (dropout 0.2)



dropout 0.2 twice




dropout 0.4 twice




Other experiments

- ❖ We experimented with many different layer combinations and even some pretrained models
- ❖ All experiments are compared with the 'final model' described previously
- ❖ Two most relevant attempts: modifying the sequence (using the same layers but in different order), adding significantly more layers


default



modified sequence



default



(note: different scale)

significantly more layers




Optimizer


- ❖ Specified during model compilation
- ❖ Specific algorithms used for model training
- ❖ Have different speed and performance based on data and other parameters
- ❖ Different parameters: learning rate, decay rate, functions
- ❖ Affect loss function – try to reduce losses

```
model.compile(optimizer='nadam',  
              loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),  
              metrics=['accuracy'])
```


Adam



Adadelata



NAdam




Epochs

- ❖ Amount of training iterations
- ❖ Each iteration, all samples are passed through the neural network once
- ❖ More does not necessarily mean better – after a certain amount of iterations, accuracy stays nearly constant or even goes down
- ❖ For research, we mostly used epochs=50 – good middleground between short training time and decent results

```
epochs = 400
history = model.fit(
    train_ds,
    validation_data=val_ds,
    epochs=epochs
)
```

Final model performance



Conclusions

- ❖ The dataset we used was not big enough to fully prevent overfitting, and to keep training accuracy as close as possible to validation accuracy
- ❖ We had to settle for some unoptimized parameters (eg. `batch_size`) – hardware problems
- ❖ The final model performs very well on training data

Problems we encountered

- ❖ Main problem – because the model is only trained on 37 classes of data, it cannot distinguish other classes (other breeds of dogs or cats)
- ❖ When passing an image that does not belong to these 37 classes, we will not get good results

Sources

- ❖ <https://www.kaggle.com/datasets/zippyz/cats-and-dogs-breeds-classification-oxford-dataset>
- ❖ <https://www.tensorflow.org/>
- ❖ <https://docs.python.org/3/reference/>
- ❖ <https://machinelearningknowledge.ai/>
- ❖ <https://towardsdatascience.com/>
- ❖ <https://anhreynolds.com/blogs/cnn.html>
- ❖ <https://iq.opengenus.org/purpose-of-different-layers-in-ml/>